

# CPSC 457

## Processes - part 1

Contains slides from Mea Wang, Andrew Tanenbaum and Herbert Bos

# Multitasking

- multitasking = concurrent execution of multiple programs
- allows computer to run **N** processes with **M** CPUs, even if **N** > **M**
- works even with a single CPU:

*repeat forever - using a timer interrupt*  
*give CPU to program (i) for a short time*  
*give CPU to program (i+1) for a short time*  
*give CPU to program (i+2) for a short time*  
*...*

- just an illusion of parallelism – programs do not have to execute their instructions exactly at the same time, as long as it appears that way
- just like multiprogramming:
  - multitasking allows us to reduce CPU idling during I/O,  
CPU can be given to another program rather than remain idle
  - multitasking is only practical when memory is big enough to hold multiple running programs

# Multiprogramming and multitasking

- early computers had limited memory
- only one program could run at a time
- lengthy I/O → idle CPU

- cheaper memory → multiple programs can be loaded simultaneously
- **multiprogramming** – OS gives CPU to another program if current program must wait on I/O

- **cooperative multitasking** - programs can voluntarily yield CPU to another program
- early Windows and Mac OS implemented this
- today you can still use `sleep(0)` or `sched_yield()`

- **preemptive multitasking** – a program gets a fraction of a second to execute, then OS automatically switches to the next program, and so on
- nearly all modern OSes implement this

- **multithreading** – allows even more efficient multitasking, usually preemptive

# Process

---

- modern OSes run multiple programs at the same time (multitasking)
- OS needs to keep track and of all running programs
- OS needs to prevent accidental or malicious interference between programs
- OS maintains some information about each **running program** – we call this a **process**


# Program != Process

---

- a program is a *passive* entity – executable file containing a list of instructions, eg. stored on disk
- a process is an *active* entity – associated with a unique program counter and other resources
- a program *becomes* a process when it is loaded into memory for execution
- a single program can be used to start multiple processes
  - eg. running multiple terminals or shells

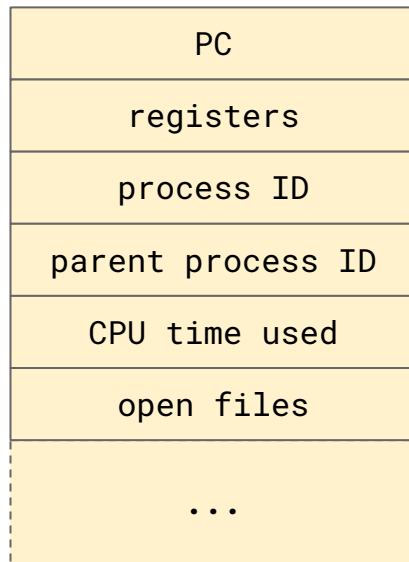
# A process in memory

- each process gets its own **address space**
  - part of memory available to a process, decided by OS
  - on modern OSes it is a *virtual* address space (0 - max), isolated from other processes
- examples of things in address space of a process:
  - **text section**: the program code
  - **data section**: global variables, constant variables
  - **heap**: memory for dynamic allocation during runtime
  - **stack**: temporary data (parameters, return address, local variables)
  - plus many other bits of information needed by the OS for management
- OSes often group all information needed to run a process in a data-structure that we will call **Process Control Block**



# Process control block (PCB)

- typical parts of PCB:
  - process state
  - program counter, CPU registers
  - priority, pointers to various queues
  - memory management info: eg. page tables, segment tables, etc.
  - accounting info: eg. CPU time, timeout values, process numbers, etc.
  - I/O status info: open files, I/O devices, etc.
- PCB in practice is not a single data structure
- a **process table** is a collection of all PCBs



# More examples of fields of a PCB

## Process management

program counter  
registers  
stack pointer  
process state  
priority  
scheduling parameters  
process ID  
parent process  
process group  
signals  
process start time  
CPU time used  
children's CPU time used  
time of next alarm  
...

## Memory management

pointer to text segment  
pointer to data segment  
pointer to stack segment  
...

## File management

root directory  
working directory  
file descriptors  
user ID  
group ID  
...

Look for “**task\_struct**” in Linux kernel sources

<https://github.com/torvalds/linux/blob/master/include/linux/sched.h>



# Operations on processes

- processes need to be created and deleted dynamically
- any multitasking OS must provide mechanisms (APIs) for this
  - in UNIX *process creation* is accomplished by using `fork()`
    - **parent process** – the process that is creating a new process calls `fork()`
    - **child process** – the newly created process
    - processes in the system form a **process tree**
    - each process gets **PID** - a unique process identifier, usually an `unsigned int`
  - *process execution*, eg. `fork()` in Unix
  - *process termination*, eg. `exit()` or `kill()`
    - to let the OS know it can delete the process and free up resources
    - termination can be (typically) only requested by the process itself, its parent, or an unrelated process provided it's owner has the right permissions
  - many other operations exist as well, such as synchronization, communication, ...

# A multiprocess program in C

```
$ man fork
```

```
pid_t fork(void);
```

`fork()` creates a new process by duplicating the calling process. The new process is referred to as the child process. The calling process is referred to as the parent process.

The child process and the parent process run in separate memory spaces. At the time of `fork()` both memory spaces have the same content. Memory writes, file mappings (`mmap(2)`), and unmappings (`munmap(2)`) performed by one of the processes do not affect the other.

The child process is an exact duplicate of the parent process except for the following points:

...

# A multiprocess program in C

```
#include <stdio.h>
#include <unistd.h>

int main()
{
    printf("Hello\n");
    /* create & run child process - a duplicate of parent */
    fork();
    /* both parent and child will execute the next line */
    printf("world.\n");
}
```

Possible output:

???

# A multiprocess program in C

```
#include <stdio.h>
#include <unistd.h>

int main()
{
    printf("Hello\n");
    /* create & run child process - a duplicate of parent */
    fork();
    /* both parent and child will execute the next line */
    printf("world.\n");
}
```

Possible output:

Hello  
world.  
world.

Are other outputs possible?

# A multiprocess program in C

```
#include <stdio.h>
#include <unistd.h>
```

<https://repl.it/@pfederl/fork-hello-world>  
<https://repl.it/@pfederl/fork-hello-world-one-char-at-a-time>

```
int main()
{
    printf("Hello\n");
    /* create & run child process - a duplicate of parent */
    fork();
    /* both parent and child will execute the next line */
    printf("world.\n");
}
```

Possible output:

Hello  
world.  
world.

Another possible output:

Hello  
worworld.  
rld.

assuming printf()  
is unbuffered and  
outputs 1 character  
at a time

Are other outputs possible?

# A multiprocess program in C

```
#include <stdio.h>
#include <unistd.h>

int main()
{
    printf("Hello\n");
    /* create & run child process - a duplicate of parent */
    fork();
    /* both parent and child will execute the next line */
    printf("world.\n");
}
```

If `fork()` fails, it returns `-1`. You should always check the return value of a system call.

Possible output:

```
Hello
world.
world.
```

Another possible output:

```
Hello
worworld.
rld.
```

Another possible output:

```
Hello
world.
```

# A multiprocess program in C

```
int main()
{
    /* create & run child process - a duplicate of parent
     * and remember the return value */
    pid_t pid = fork();
    /* both parent and child will execute the next line,
     * but will have different value for pid:
     * 0 for child, positive integer for parent, or -1 for error */
    printf("fork returned %d.\n", pid);
}
```

Possible output:

fork returned 7.  
fork returned 0.

Possible output:

fork returned 0.  
fork returned 7198.

Possible output:

fork returned -1.

# Exercise – to stay in shape

## GO AHEAD

```
int main()
{
    fprintf( stderr, "A\n");
    fork();
    fprintf( stderr, "B\n");
    fork();
    fprintf( stderr, "C\n");
}
```

## PREDICT MY OUTPUT

Can you predict  
all possible  
outputs  
assuming  
`fork()` does  
not fail?




# Exercise – to stay in shape

```
int main()
{
    fprintf( stderr, "A\n");
    fork();
    fprintf( stderr, "B\n");
    fork();
    fprintf( stderr, "C\n");
}
```

Hint:

- find all topological orderings in the graph



- then remove duplicates
- hard problem in general, but quite easy for a small graph

# Another forking exercise – predict all outputs

```
int main()
{
    for(int i=0 ; i<4 ; i++ ) {
        fork();
    }
    printf("X");
}
```

Easy version:  
assume `fork()`  
**does not** fail

Hard version:  
assume `fork()`  
**could** fail

# Another forking exercise – predict all outputs

```
int main()
{
    for(int i=0 ; i<4 ; i++ ) {
        fork();
    }
    printf("X");
}
```

Hint: are  
these  
equivalent?

```
int main()
{
    fork();
    fork();
    fork();
    fork();
    printf("X");
}
```

# Another forking exercise – predict all outputs

```
int main()
{
    for(int i=0 ; i<4 ; i++ ) {
        fork();
        printf("%d",i);
    }
}
```

This is actually  
very similar to  
the first  
exercise...

# Exercise – can you predict the output?

```
int x = 10;
int main()
{
    printf("x=%d\n", x);
    fork();
    x ++;
    printf("x=%d\n", x);
}

// assume fork() does not fail
```

**Hint:**  
child has its own  
'x' variable,  
different from  
the parent

# Fork bomb



**Do not run this on CPSC servers.**

```
int main() {  
    while(1) {  
        fork();  
    }  
    printf("X"); // ???  
}
```



# Starting an external program (programmatically)

---

- how do we start an external program in Unix?
- no dedicated system call for this purpose
- we have to `fork()` a new process
- then we *replace* child process with an external program using `exec()` system call

# A multiprocess program in C

```
$ man -s 3 exec
```

```
int execl(const char *path, const char *arg, ...);  
int execlp(const char *file, const char *arg, ...);  
...
```

The `exec()` family of functions replaces the current process image with a new process image. The functions described in this manual page are front-ends for `execve(2)`. (See the manual page for `execve(2)` for further details about the replacement of the current process image.)

The initial argument for these functions is the name of a file that is to be executed.

...



# A multiprocess program in C

```
$ man execve
```

```
int execve(const char *filename, char *const argv[], char *const envp[]);
```

`execve()` executes the program pointed to by filename. filename must be either a binary executable, or a script starting with a line of the form:

```
#! interpreter [optional-arg]
```

For details of the latter case, see "Interpreter scripts" below.

argv is an array of argument strings passed to the new program. By convention, the first of these strings should contain the filename associated with the file being executed. envp is an array of strings, conventionally of the form `key=value`, which are passed as environment to the new program. Both argv and envp must be terminated by a null pointer.

...

# Starting external program with `fork()` + `exec()`

```
int main()
{
    pid_t pid = fork();
    if (pid < 0) { /* check for error */
        fprintf(stderr, "Fork failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", "-l", NULL); /* replace process with 'ls -l' */
        /* we should be checking for errors above... */
        printf("This _should_ never print...\n");
    }
    else { /* parent process will wait for the child to complete */
        printf("Waiting for child process %d\n", pid);
        while (wait(NULL) > 0); // wait(NULL); would work in _this_ case
        printf("Child finished.\n");
        exit(0);
    }
}
```

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/wait.h>
```

# A multiprocess program in C

```
#include <stdio.h>
#include <stdlib.h>

int main()
{
    printf("Before ls.\n");
    system("/bin/ls -l");
    printf("After ls.\n");
}
```

```
$ man system
```

```
...
```

The system() library function uses fork(2) to create a child process that executes the shell command specified in command using execl(3) as follows:

```
execl("/bin/sh", "sh", "-c", command,
      (char *) 0);
```

system() returns after the command has been completed.

```
...
```

# A multiprocess program in C

```
#include <stdio.h>
#include <stdlib.h>

int main()
{
    FILE * fp = popen("/bin/ls -l", "r");
    if (fp == NULL) {
        fprintf( stderr, "popen failed.\n");
        exit(-1);
    }

    char buff[4096];
    while (fgets(buff, sizeof(buff), fp) != NULL)
        printf("%s", buff);

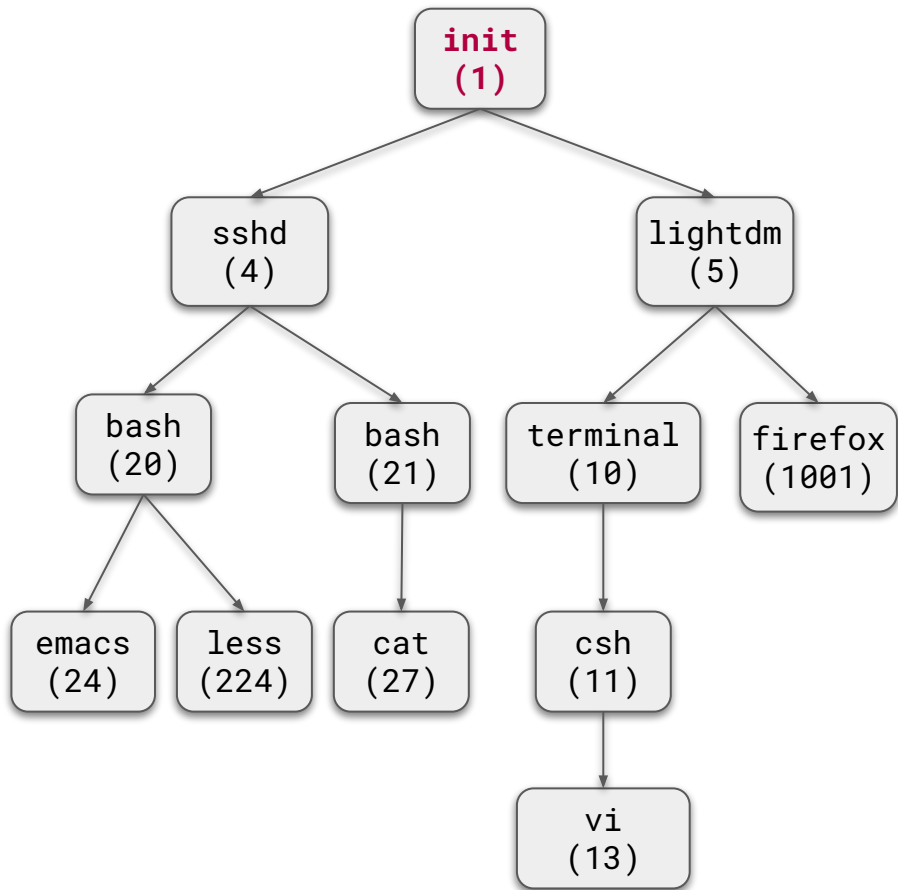
    pclose(fp);
}
```

\$ **man popen**

...

The `popen()` function opens a process by creating a pipe, forking, and invoking the shell. Since a pipe is by definition unidirectional, the type argument may specify only reading or writing, not both; the resulting stream is correspondingly read-only or write-only.

# Process tree



- **parent process** - the creating process
- **child process (child)** - the newly created process
- **PID** - the unique process identifier for each process
- in Unix, parent and child processes continue to be associated, forming a process hierarchy
- in Windows, all processes are equal, the parent process can give the control of its children to any other process

# init process

- `init` is the first process started after booting
  - older UNIX systems used init based systems
  - many newer Linux systems switched from init to `systemd`
- `init` is the `ancestor` of all user processes (direct or indirect parent), i.e. root of process tree
- `init` always has `PID = 1`
- `orphaned` processes are adopted by `init` (parent terminates before child)
- printing a process tree

```
$ pstree
$ ps axjf
```
- note: some special 'system processes' are created by kernel during bootstrap, and do not have to be descendants of `init`, such as `swapper` and `pagedaemon`

# Questions?

# Review

---

- Which one of the following executes in kernel mode?
  - ☐ A user program
  - ☐ A library function call
  - ☐ A system call
  - ☐ A system call wrapper function
  
- In C, **printf()** is a system call.
  - ☐ True
  - ☐ False



# Review

---

- When 4 programs are executing on a computer with a single CPU, how many program counters are there?
- When does a program become a process?

- What is the name of the PCB data structure in Linux?
- Name some of fields in a PCB.
- On UNIX systems, what is the name of the process that is the ancestor of all user processes?