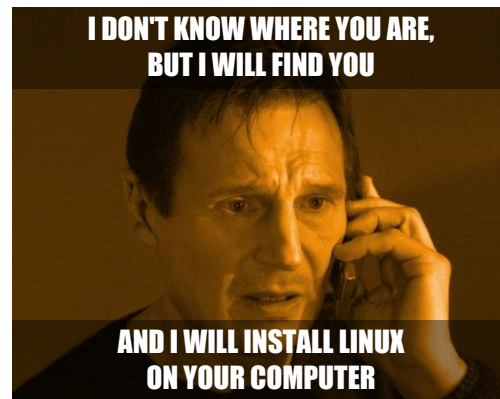


# CPSC 457



## Processes - part 2

Contains slides from Mea Wang, Andrew Tanenbaum and Herbert Bos

- CPU utilization
- processes creation, termination
- process scheduling
- process states
- context switching
- signals



# Multiprogramming on a single CPU



# CPU utilization

## ■ example:

- OS is running 4 processes, P1, P2, P3 and P4
  - P1 spends **40%** of the time waiting on I/O
  - P2 spends **20%** of the time waiting on I/O
  - P3 spends **50%** of the time waiting on I/O
  - P4 spends **90%** of the time waiting on I/O
- if there is only one CPU, what will be the CPU utilization?  
i.e. what percentage of the time is the CPU going to be running 'something'?

## ■ Answer:

- CPU utilization = probability that at least one of the processes is not waiting on I/O  
= ???

# CPU utilization

## ■ example:


- OS is running 4 processes, P1, P2, P3 and P4
  - P1 spends **40%** of the time waiting on I/O
  - P2 spends **20%** of the time waiting on I/O
  - P3 spends **50%** of the time waiting on I/O
  - P4 spends **90%** of the time waiting on I/O
- if there is only one CPU, what will be the CPU utilization?  
i.e. what percentage of the time is the CPU going to be running 'something'?

## ■ Answer:

- CPU utilization = probability that at least one of the processes is not waiting on I/O
  - =  $1 - (\text{probability that all processes are waiting on I/O})$
  - =  $1 - (0.4 * 0.2 * 0.5 * 0.9) = 0.964 = 96.4\%$

# CPU utilization - under simplistic multiprogramming model

- assume  $n$  similar processes
- each process spends the same fraction  $p$  of its time waiting on I/O
- then CPU utilization =  $1 - p^n$



CPU utilization as a function of the number of processes in memory.

# CPU utilization example

---

- example:
  - computer has 8GB of RAM
  - 2GB are taken up by OS, leaving 6GB available to user programs
  - user wants to run multiple copies of a program that needs 2GB RAM, with average 80% I/O
  - with 6GB remaining, user could run 3 copies of the program
  - CPU utilization would be  $= 1 - 0.8^3 \approx 49\%$
- is it a good idea to buy 8GB more of RAM?

# CPU utilization example

## ■ example:

- computer has 8GB of RAM
- 2GB are taken up by OS, leaving 6GB available to user programs
- user wants to run multiple copies of a program that needs 2GB RAM, with average 80% I/O
- with 6GB remaining, user could run 3 copies of the program
- CPU utilization would be  $= 1 - 0.8^3 \approx 49\%$

## ■ is it a good idea to buy 8GB more of RAM?

- with 14GB available, we could run 7 copies of the program
- CPU utilization would be  $= 1 - 0.8^7 \approx 79\%$
- throughput increased by  $79\% - 49\% = \mathbf{30\%}$

## ■ is it a good idea to buy 8GB more?

- we could run 11 programs  $\rightarrow$  CPU utilization  $= 1 - 0.8^{11} \approx 91\%$
- throughput increased only by  $91\% - 79\% = \mathbf{12\%}$  (diminishing returns)



# Process creation

---

- in UNIX
  - `init` process is created at boot time by kernel (special case)
  - afterwards, only an existing process can create a new processes, via `fork()`
  - therefore all other processes are descendants of `init`
  - in many modern Linux distributions `init` is replaced by `systemd`
  - `fork()` often followed by `exec*()` to spawn a different program  
(remember, you you can use `system()` and `popen()` convenience functions)
- in Windows: `CreateProcess()` is used to create processes,  
but the behavior is quite different from `fork()`

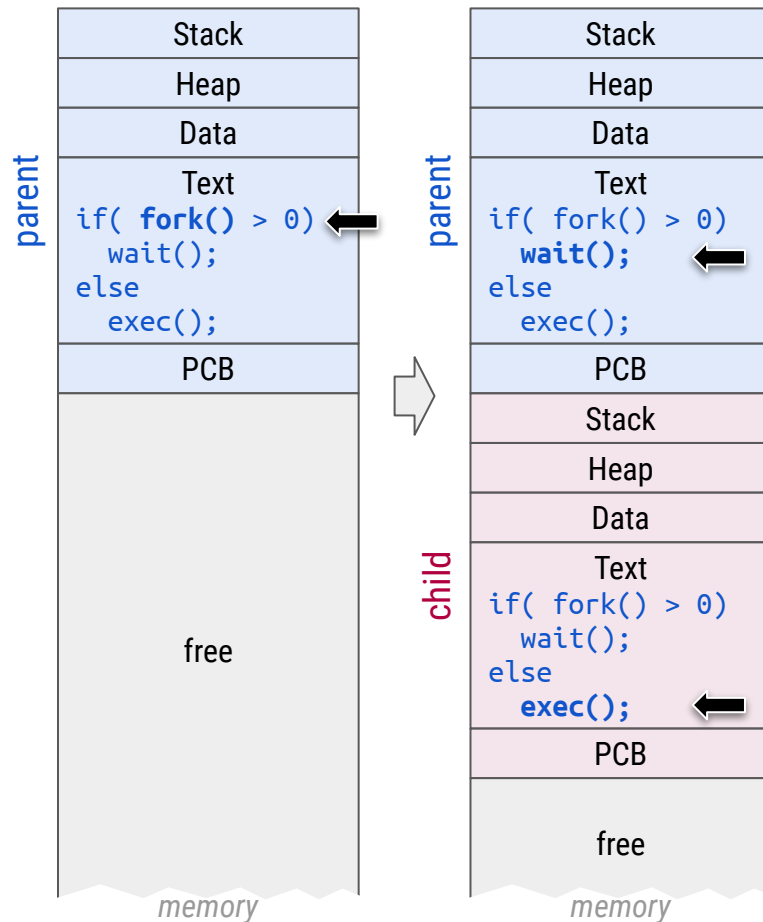
# Process creation

---

- there are many reasons for a process to create new process[es]
- during system initialization (boot)
  - spawning background processes — daemons, services, eg. database server
- application decides to spawn additional processes
  - eg. to execute external programs or to do parallel work
- a user requests to create a new process
  - eg. GUI windows manager application allows launching new applications
- starting batch jobs
  - mainframes

# Address space

- each process has its own address space
- `fork()` duplicates address space, creating nearly identical copy of itself
- next instruction is the same (**if**)
- but code flow may differ for child vs. parent



# Resource allocation

---

- several options for allocating resources for a new process, for example:
  - child obtains resources directly from the OS
    - most common, easiest to implement
    - every new process gets the same resources
    - fork bomb crashes the system
  - child obtains subset of parent's resources
    - parent must give some of its resources to child
    - fork bomb has limited impact
  - parent shares resources with the child – eg. with threads
  - hybrids

# Common parent-child execution scenarios

- after child is created, parent usually does one of 3 things:
  1. the parent waits until the child process is finished
    - often used when child executes another program, eg. `fork/exec()`, or `system()`
  2. the parent continues to execute concurrently and independently of the child process
    - eg. autosave feature
  3. the parent continues to execute concurrently, but synchronizes with the child
    - can be quite complicated to synchronize

```
pid = fork()
if pid > 0 :
    wait()
```

```
pid = fork()
if pid > 0 :
    do_whatever()
    exit()
```

```
pid = fork()
if pid > 0 :
    do_something_1()
    synchronize()
    do_something_2()
    synchronize()
    ...
```

# Process termination

- typical reasons for terminating a process
- voluntary:
  - **normal exit** - eg. application decides to terminate, or user instructs an app to 'close'
    - app calls `exit(0)` or returns `0` from `main()` – which is the same thing in C
  - **error exit** - application detects an error, optionally notifies user, and then terminates
    - app calls `exit(N)` or returns `N` from `main()` with `N!=0`
- involuntary:
  - **fatal error** – aka bugs in software
    - error detected by OS, eg. accessing invalid memory, division by zero
  - **external** – killed by another process
    - parent, or another process calls `kill()`
    - eg. during shutdown, pressing `<ctrl-c>` in terminal, closing GUI window

# Process termination

---

- parent may terminate its children for different reasons, for example:
  - the task assigned to the child is no longer required
  - the parent needs/wants to exit and wants to clean up first
- in Unix, when a parent process is terminated:
  - the child processes may be terminated, or assigned to the grandparent process, or to the `init` process
  - process hierarchy is always maintained
- default behavior on Linux is to **reparent** the child process to the `init` process
  - this can be changed (eg. to kill children, reparent to some other process)
  - see `$ man prctl` for more details

# Process termination

---

- when terminating a process the OS must clean up:
  - free memory used by the process
  - delete PCB
  - delete process from process table
  - kill children or assign them a new parent
  - close open files
  - close network connections
  - ...

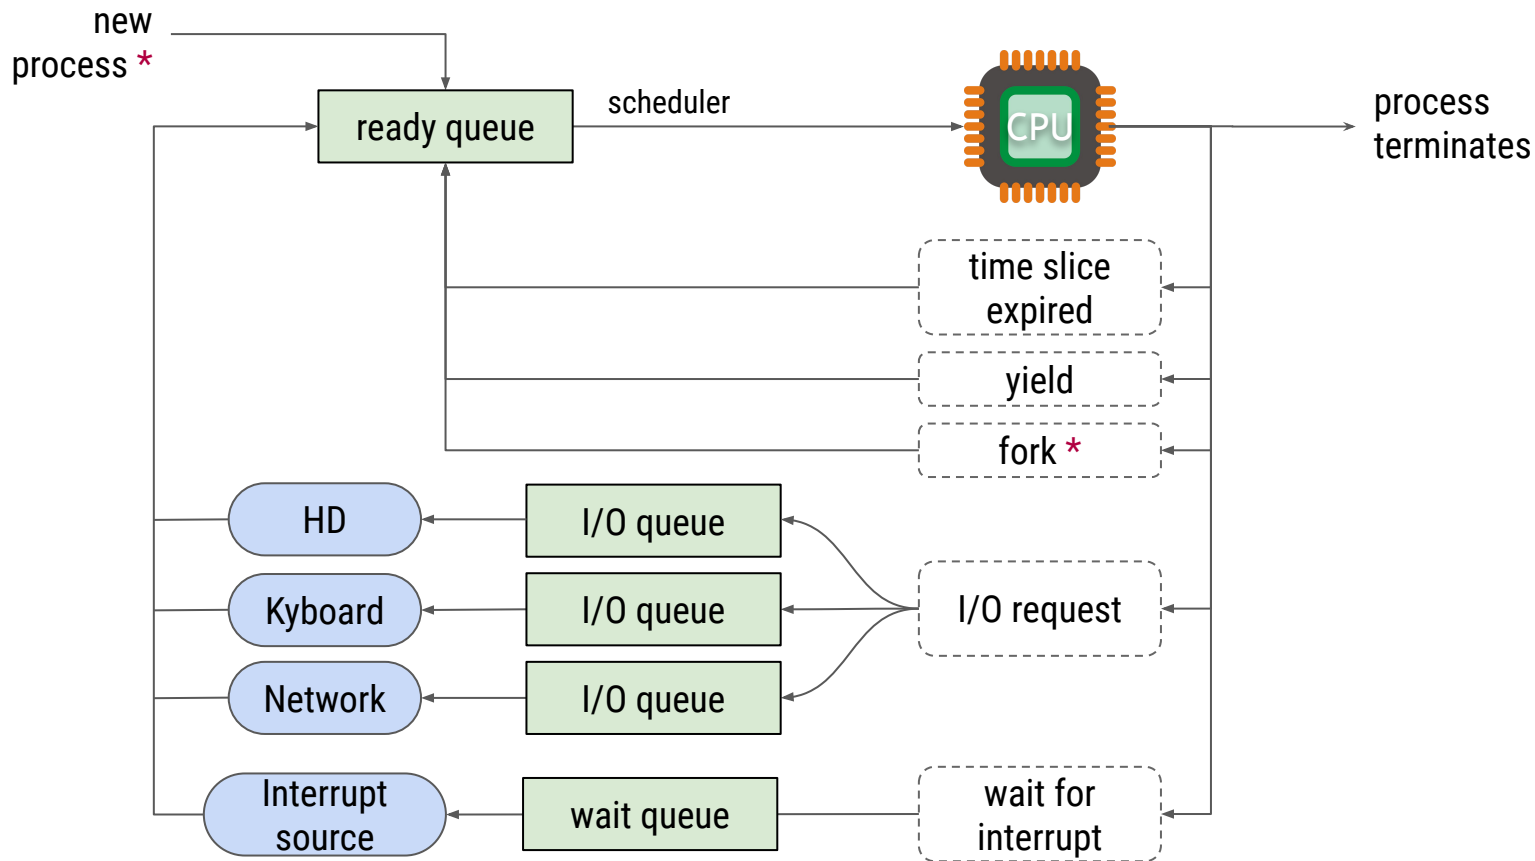


# Process scheduling

---

- part of multitasking is deciding which process gets the CPU next
- typical objective is to maximize CPU utilization
- **process scheduler:**
  - kernel routine/algorithm that chooses one of available process to execute next on the CPU
  - selected from processes in a ready queue
- OS maintains different scheduling queues:
  - job queue: all programs waiting to run, usually found in batch systems
    - eg. priority queue
  - ready queue: all processes that are ready to execute their next instruction
    - eg. linked list, priority queue, ... depends on scheduler
  - device queues: processes waiting for a particular device
    - each device has its own queue

# Process scheduling diagram



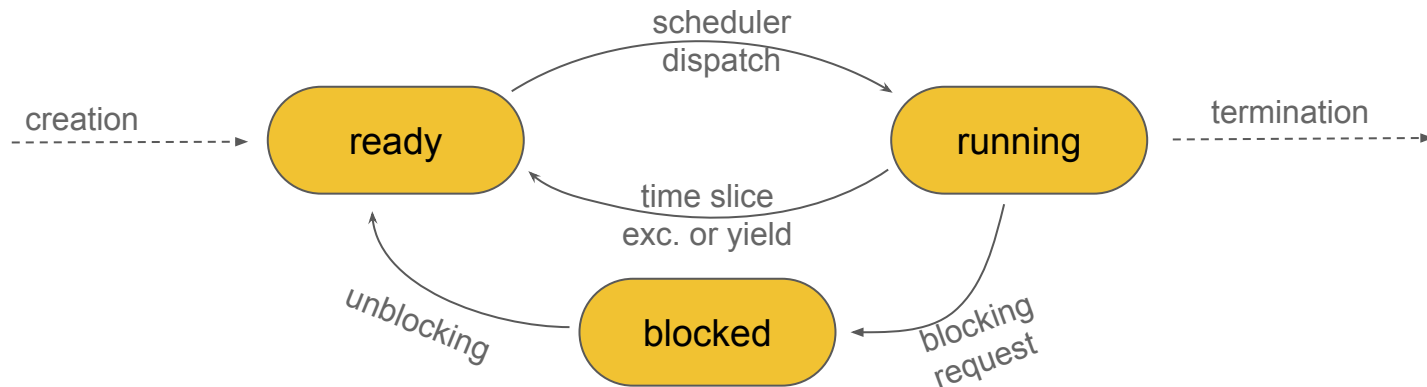
# Process states

3 process states:

- **running** — actually running on the CPU
- **blocked** — waiting for some event to occur, eg. I/O
- **ready** — the process is ready to execute on CPU

only 4 transitions are possible:

- ready → running
- running → ready
- running → blocked
- blocked → ready

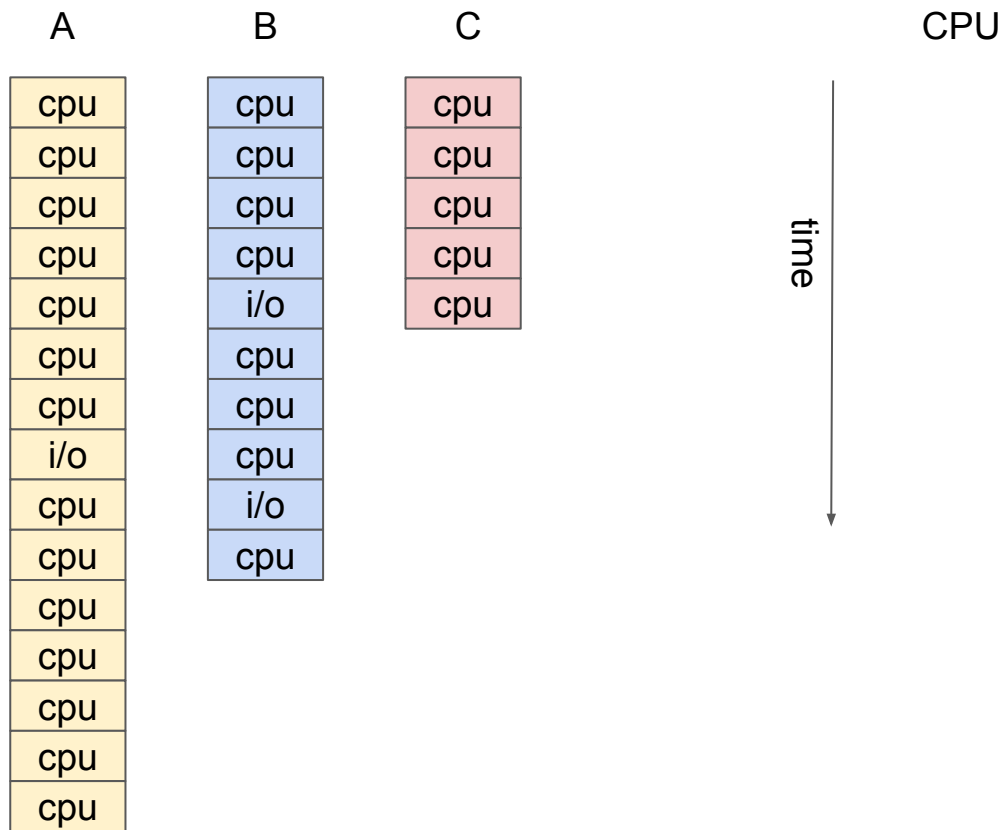


# Exercise – simulating round-robin scheduling

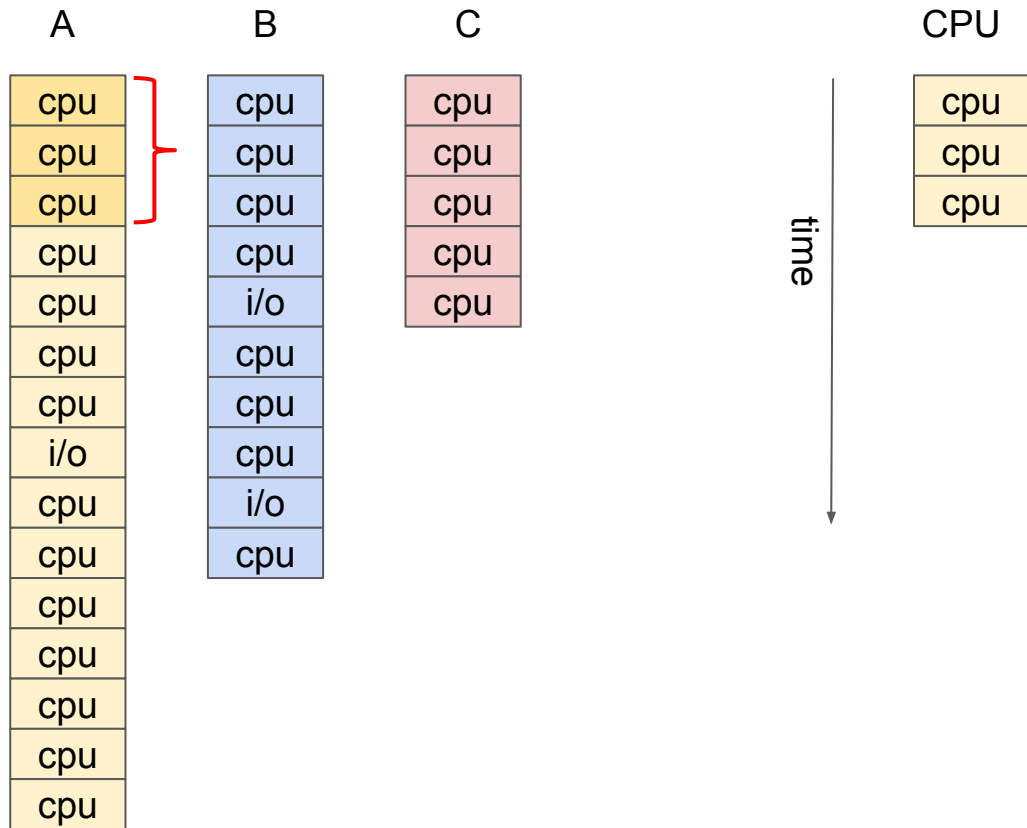
- simulate 3 processes A, B, C
  - A: 7 units of CPU, 1 unit of I/O, 7 units of CPU
  - B: 4 CPU, 1 I/O, 3 CPU, 1 I/O, 1 CPU
  - C: 5 CPU
- assume time slice of 3 units
  - each process gets 3 units of CPU cycles
  - if process requests I/O during its time slice, OS switches to the next process
  - otherwise, after time slices expires, OS switches to next process
- assume I/O is very short, less than 1 time-slice

A	B	C
cpu	cpu	cpu
cpu	cpu	cpu
cpu	cpu	cpu
cpu	cpu	cpu
cpu	i/o	cpu
cpu	cpu	
cpu	cpu	
i/o	cpu	
cpu	i/o	
cpu	cpu	
cpu		
cpu		
cpu		
cpu		
cpu		

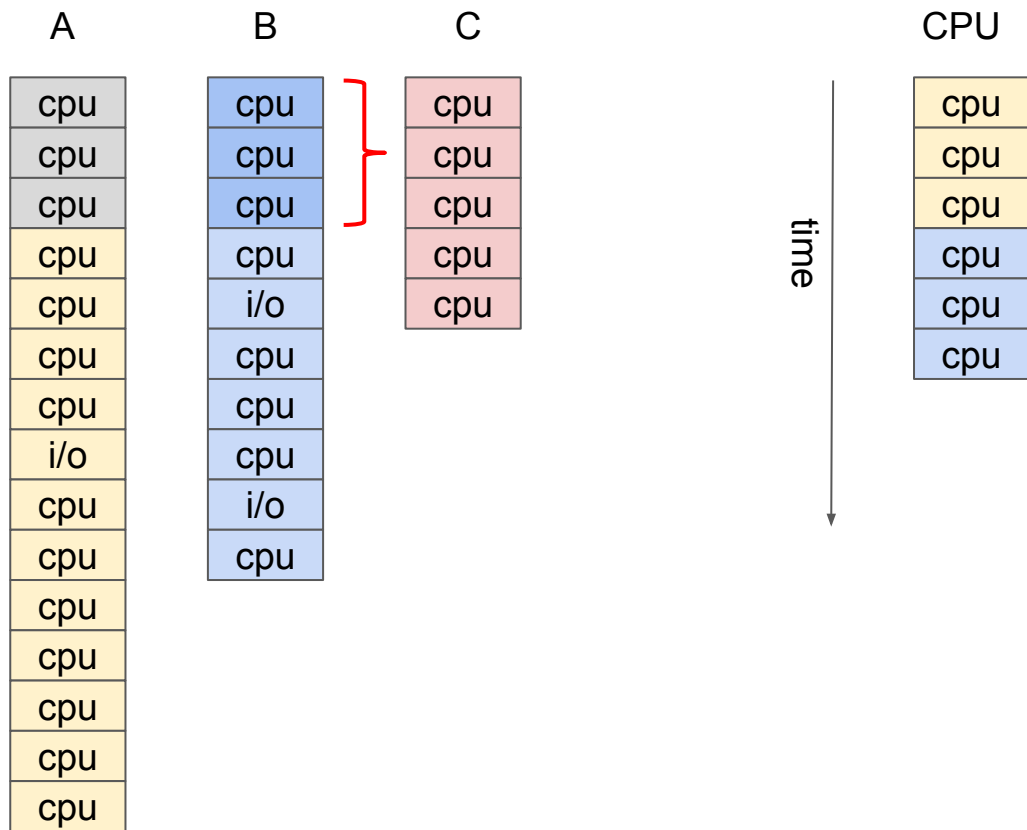
# Exercise



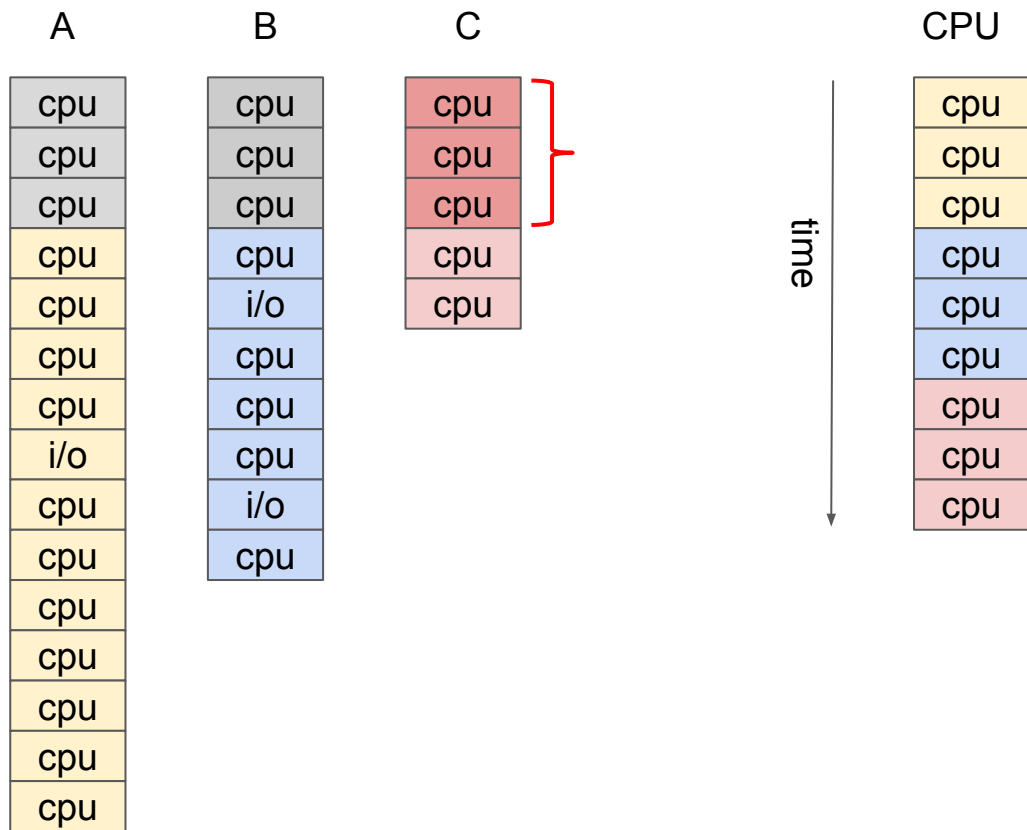
# Exercise



# Exercise

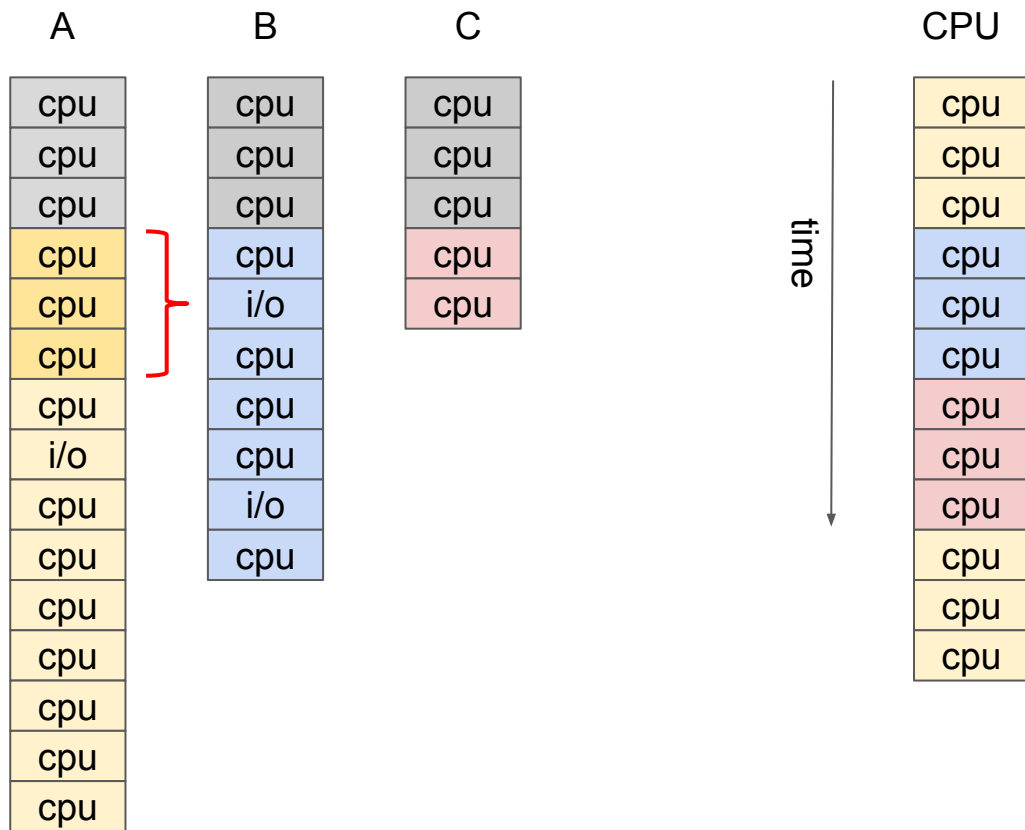



# Exercise



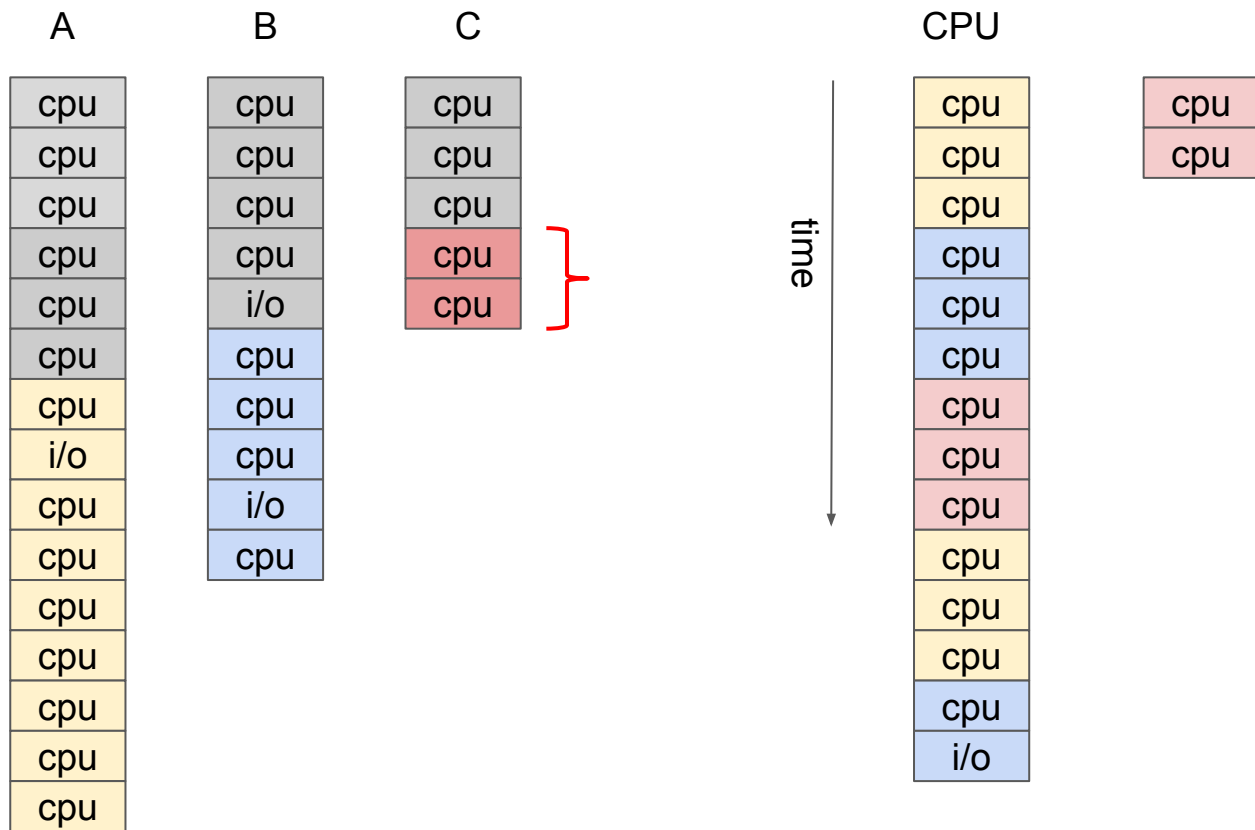


# Exercise

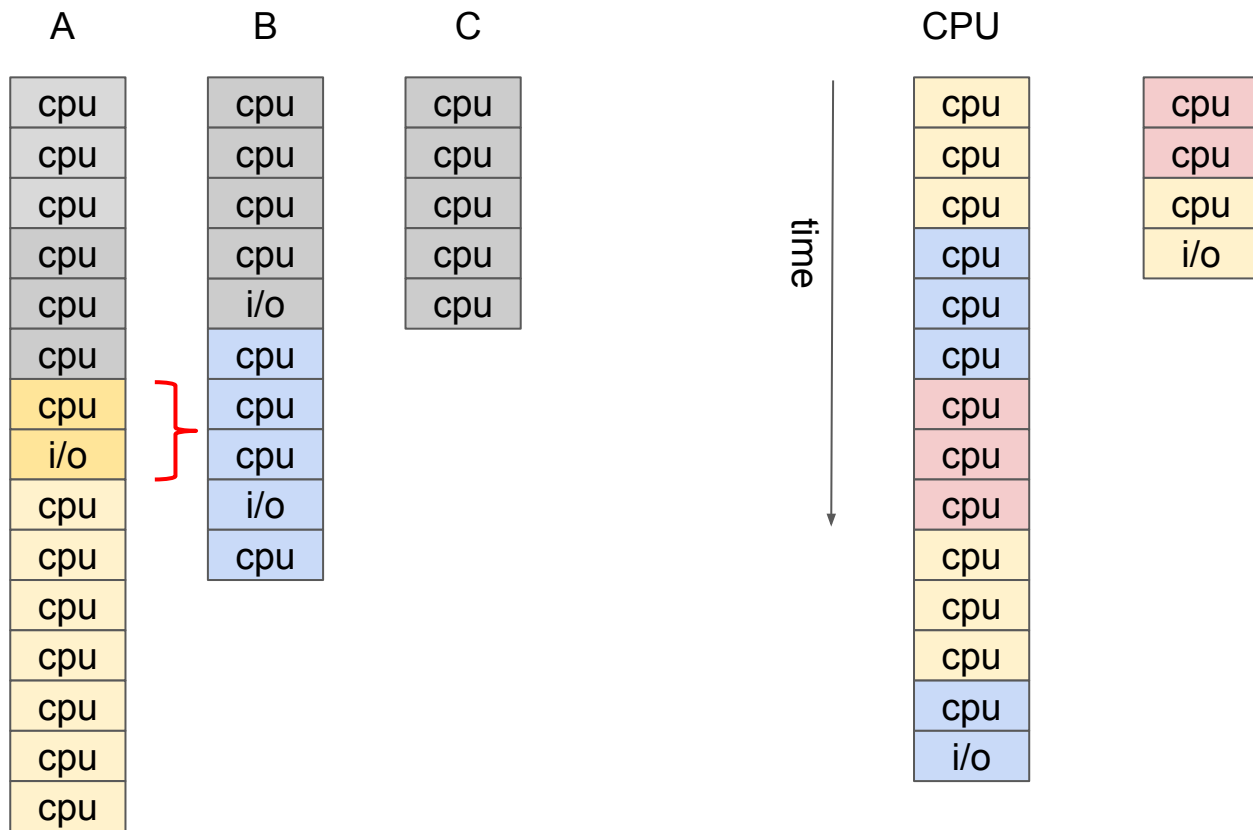





# Exercise



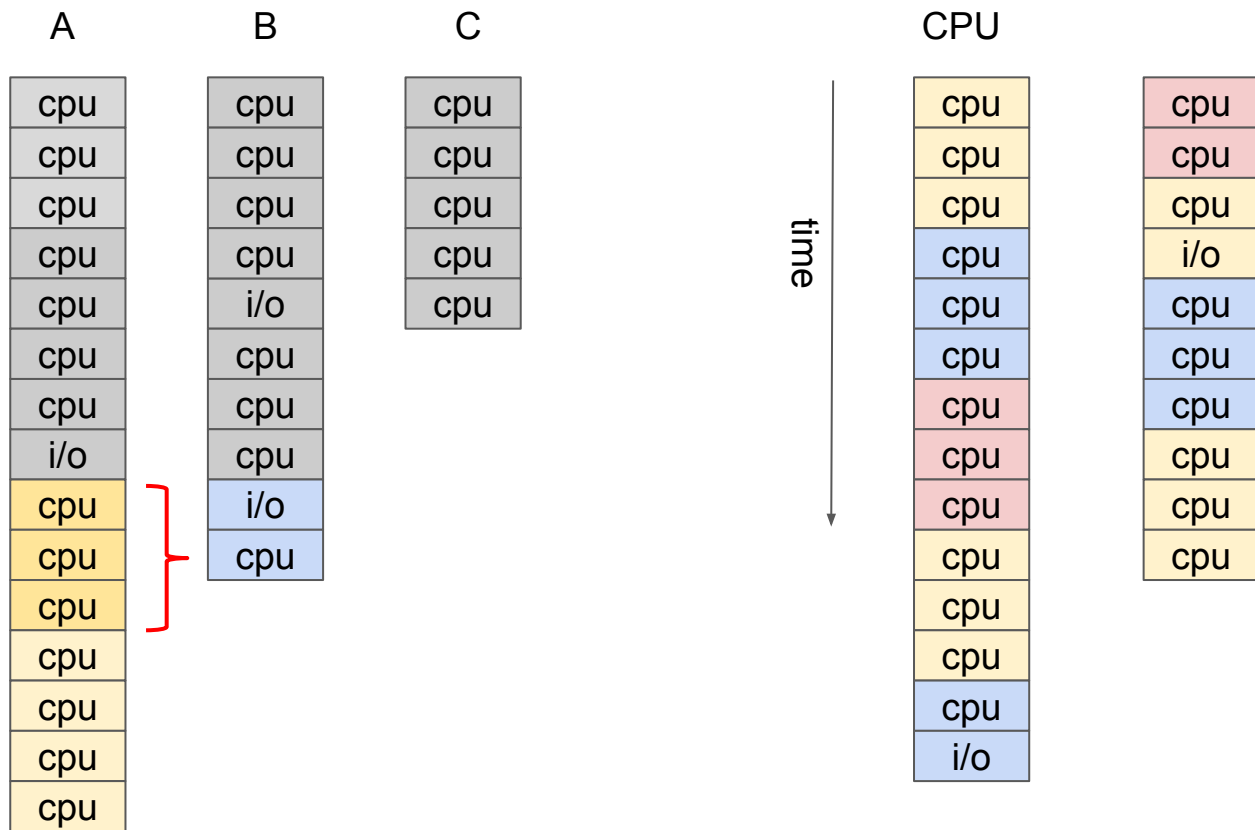
# Exercise



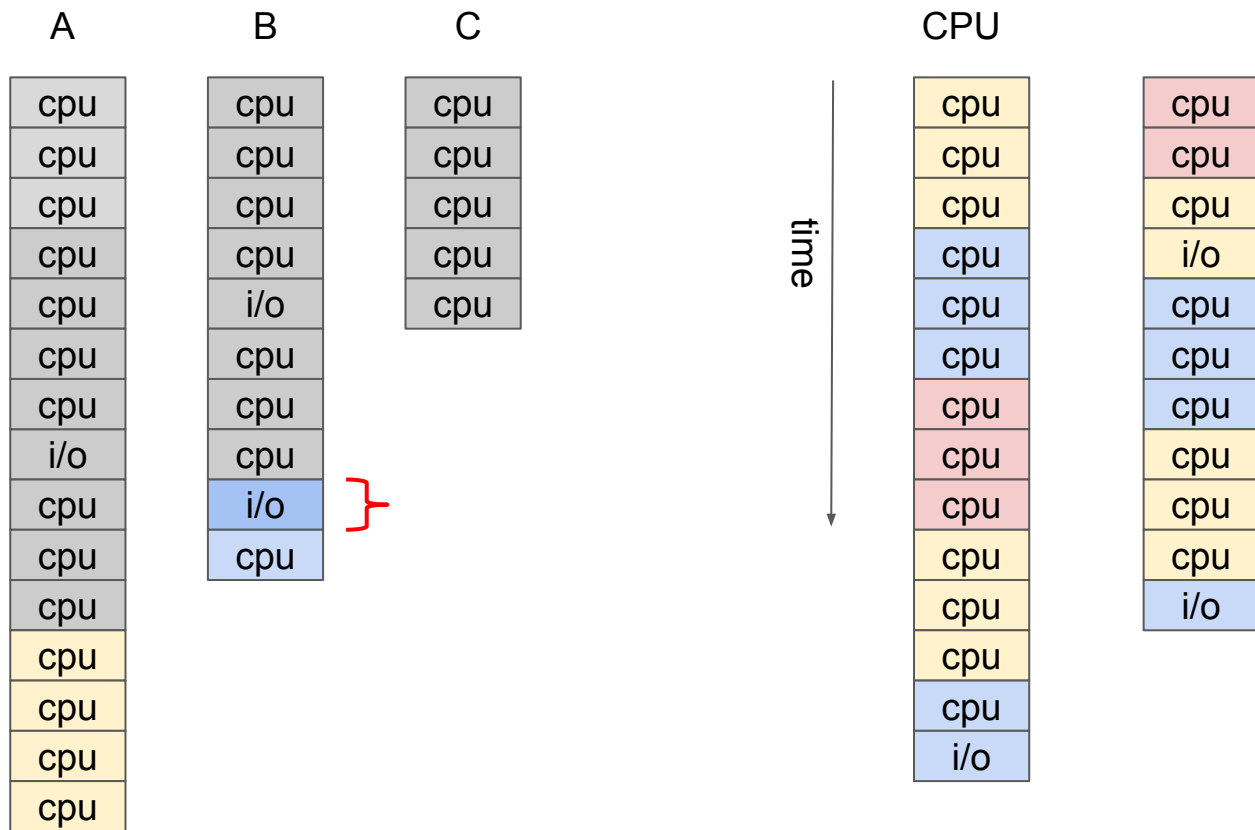
# Exercise



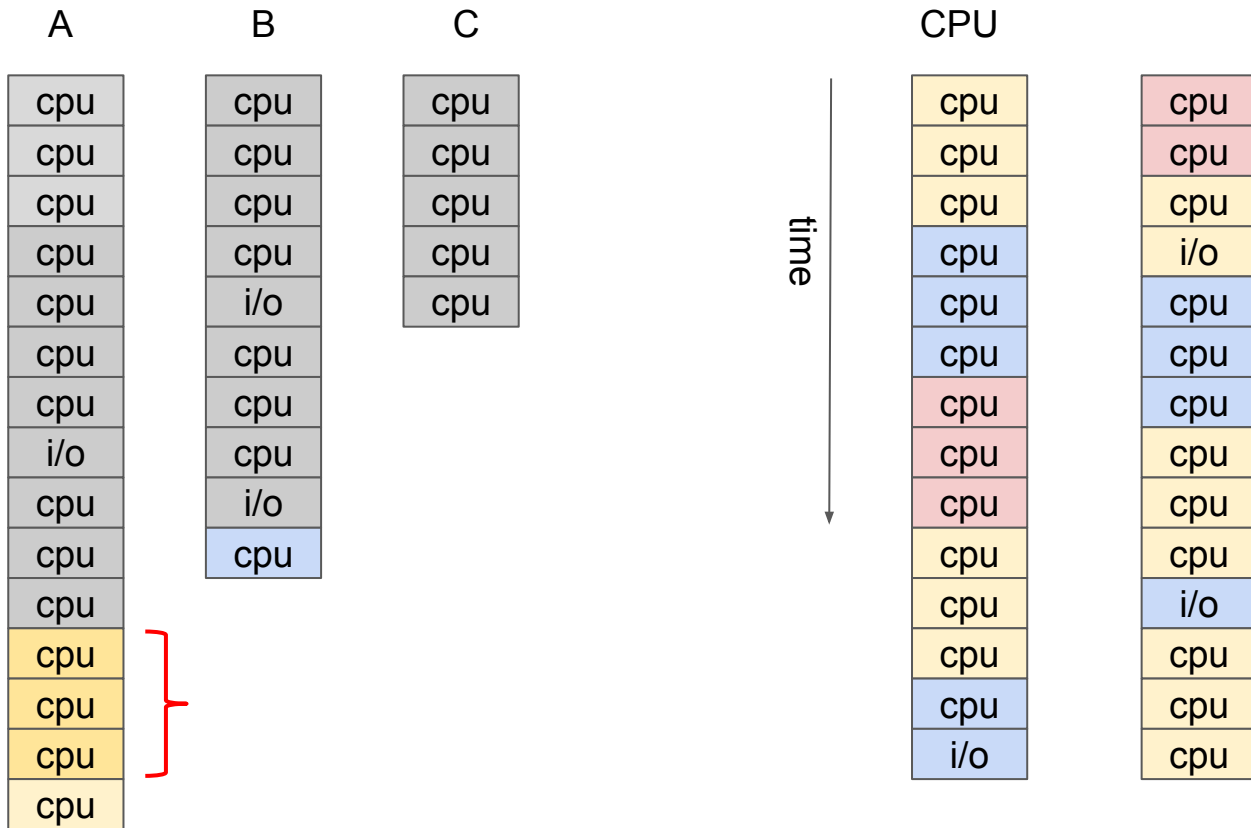
# Exercise



# Exercise




# Exercise

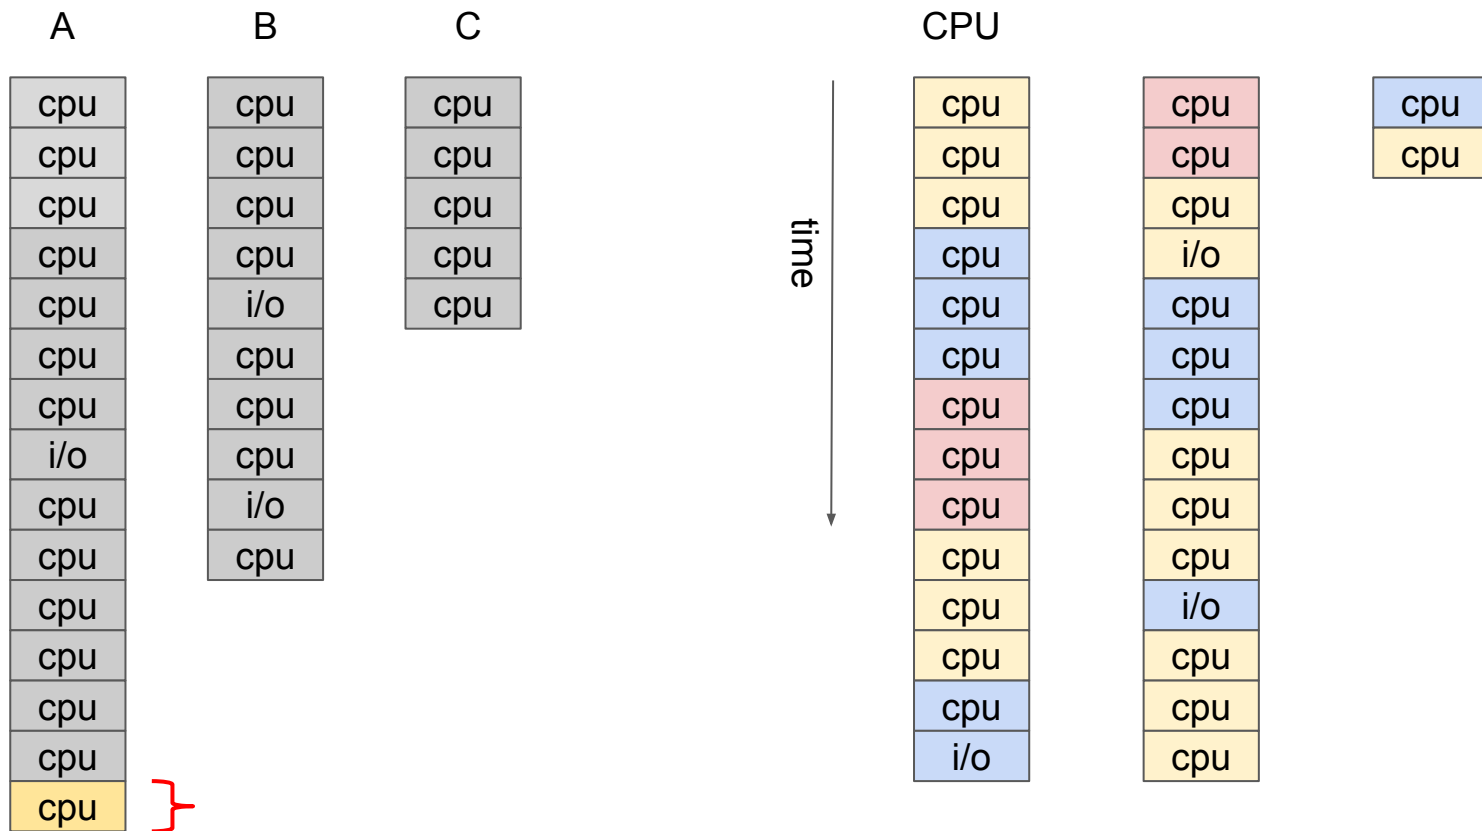




# Exercise



# Exercise



# Context switch

---

- context switch is an essential part of any multitasking OS
- we need context switching to implement multitasking when  $\# \text{ CPUs} < \# \text{ processes}$
- allows (the illusion of) sharing of a small number of CPUs among many processes
  
- OS maintains a context (state) for each process
  - usually part of PCB, includes saved registers, open files, ...
- when OS switches between processes A and B:
  - OS first saves A's state in A's PCB
    - eg. save current CPU registers into  $\text{PCB}_A$
  - OS then restores B's state from B's PCB
    - eg. load CPU registers from  $\text{PCB}_B$

# Context switch

- context switch occurs in kernel mode:
  - for example when process exceeds its **time slice**
    - enforcing time slice policy usually implemented via timer interrupt
  - or when current process voluntarily relinquishes (**yields**) CPU, eg. by sleeping
  - or when current process requests a blocking I/O operation, or any blocking system call
  - or due to other events, such as keyboard, mouse, network interrupts
- context switch introduces time overhead
  - CPU spends cycles on no "useful" work, eg. saving/restoring CPU registers
  - context switch routine is one of the most optimized parts of kernels
- context switch performance could be improved by hardware support:
  - eg. some CPUs support saving/restoring registers in a single instruction
  - or CPU could support multiple sets of registers
  - software based context switch is slower, but more customizable, and often more efficient

# Unix signals

- a form of **interprocess communication (IPC)**
- similar concept to hardware interrupts on CPUs (you can think of it as process interrupt)
- signals are asynchronous
- very limited form of IPC - the message is a single number, from a set of predefined integers

```
$ man -s 7 signal
```

```
...
```

Signal	Value	Action	Comment
SIGHUP	1	Term	Hangup detected on controlling terminal or death of controlling process
SIGINT	2	Term	Interrupt from keyboard
SIGQUIT	3	Core	Quit from keyboard
SIGILL	4	Core	Illegal Instruction
SIGABRT	6	Core	Abort signal from abort(3)
SIGFPE	8	Core	Floating-point exception
SIGKILL	9	Term	Kill signal
SIGSEGV	11	Core	Invalid memory reference
SIGPIPE	13	Term	Broken pipe: write to pipe with no readers; see pipe(7)

# Unix signals

---

- signals are used to notify a process that a particular event has occurred
  - one process (or thread) sends a signal, another process (or thread) receives it
  - it is possible for a process to signal itself
  - kernel can send signals to any processes
- signal lifetime:
  - a signal is **generated/sent**, usually as a consequence of some event
  - signal is **delivered/pending** to a process
  - delivered signal is **handled** by the process via **signal handler**
  - some signals can be **ignored** - signal delivered to a process that ignores it is lost
  - some signals can be **blocked** - signal stays pending until it is unblocked

# Generating signals

- manually from one process to another process

```
kill( pid, signal); // pid can be the current process
```

- periodically via timer

```
alarm() or setitimer()
```

- by kernel — to handle exceptions

eg. on segmentation fault kernel sends **SIGSEGV**

```
int main() {  
    * (char *) 0 = 1;  
}
```

- from command line

```
$ kill 12345    # tries to kill a specific process with signal SIGTERM
```

```
$ kill -9 12345 # kills a specific process with signal 9 → SIGKILL
```

```
$ kill -9 -1    # kills all processes except pid=1 (init)
```

- more information on signals

```
$ man -s 7 signal
```

# Signal handling

- **signal handler** - a function that will be invoked when a signal is delivered
- **default signal handler** - all programs start with default handlers with default behaviours
- **a user-defined signal handler** - programs can override the default handlers
  - signals handled by a user-defined handler are '**caught signals**'
- not all signals can be **caught**
  - `$ kill -9 pid` always kills the process because `SIGKILL(9)` cannot be caught, blocked or ignored, and default handler kills the process
  - `<ctrl-c>` in a terminal will deliver `SIGINT(2)` to the running process, which can be caught, ignored or blocked
  - `<ctrl-z>` in a terminal will deliver `SIGSTOP(2)` signal to the running process, which cannot be caught, ignored or blocked, default handler suspends the process



# Signal handling

- signals can be delivered anytime, even while your program is in the middle of a function, or in the middle of applying an operator (C++)
  - the state of your data might be in an inconsistent state
  - also signal handler could itself be interrupted !!!
- when writing a signal handler:
  - keep it simple, for example:
  - modify a global flag and let the program handle the interrupt later
  - declare global variables with volatile keyword (eg. `volatile sig_atomic_t sigStatus = 0;`)
  - only call reentrant functions inside the handler
  - more information and advanced tips, such as preventing signals from interrupting signal handlers:

[https://www.gnu.org/software/libc/manual/html\\_node/Signal-Handling.html](https://www.gnu.org/software/libc/manual/html_node/Signal-Handling.html)

# Reentrant functions

---

- reentrant functions are functions:
  - that can be interrupted in the middle of an operation
  - and then called again (re-entered)
  - and finally the original function call can finish executing
- used in interrupt handlers, signal handlers, multi-threaded applications, recursion \*
- when writing reentrant functions:
  - don't use global variables
    - there are some exceptions, eg. using atomic operations
  - do not call non-reentrant functions
    - unless you can temporarily disable interrupts / signals

# Reentrant functions

- example of a non-reentrant function:

```
int t;  
  
void bad_swap(int *x, int *y)  
{  
    t = *x; // using a global variable t !!!  
    *x = *y;  
    // hardware interrupt or signal might  
    // result in invoking/re-entering swap() here  
    *y = t;  
}
```

- easy to fix... can you guess how?

# Reentrant functions

- example of a reentrant function:

```
void swap(int *x, int *y)
{
    int t = *x; // using a local variable t
    *x = *y;
    // hardware interrupt / signal here would be safe to call swap() again
    *y = t;
}
```

- by using a local variable, the swap() function can be interrupted and re-entered anywhere
- please note that reentrant functions, like the one above, are often also thread-safe, but not always  
see [https://en.wikipedia.org/wiki/Reentrancy\\_\(computing\)](https://en.wikipedia.org/wiki/Reentrancy_(computing)) for examples

# Signal handling example

```
#include <stdio.h> <stdlib.h> <signal.h> <unistd.h>

void sigint_handler( int signum )
{
    printf("\ncaught signal=%d\n", signum);
    printf("LOL, you think <ctrl-c> will stop me?!?!?\n");
}

int main (int argc, char *argv[])
{
    /* catch <ctrl-c> and laugh at the user */
    signal(SIGINT, sigint_handler);

    for(int i = 1 ; i < 10 ; i++) {
        printf("Loop=%d\n", i);
        sleep( 1 );
    }
    printf("Exiting now.\n");
    exit(0);
}
```

## Output:

```
$ ./a.out
Loop=1
Loop=2
Loop=3
Loop=4
Loop=5
Loop=6
Loop=7
Loop=8
Loop=9
Exiting now.
```

# Signal handling example

```
#include <stdio.h> <stdlib.h> <signal.h> <unistd.h>
```

```
void sigint_handler( int signum )
{
    printf("\ncaught signal=%d\n", signum);
    printf("LOL, you think <ctrl-c> will stop me?!?!?\n");
}
```

```
int main (int argc, char *argv[])
{
    /* catch <ctrl-c> and laugh at the user */
    signal(SIGINT, sigint_handler);

    for(int i = 1 ; i < 10 ; i++) {
        printf("Loop=%d\n", i);
        sleep( 1 );
    }
    printf("Exiting now.\n");
    exit(0);
}
```

## Possible output:

```
$ ./a.out
Loop=1
Loop=2
Loop=3
^C
caught signal=2
LOL, you think <ctrl-c> will stop me?!?!?
Loop=4
Loop=5
Loop=6
^C
caught signal=2
LOL, you think <ctrl-c> will stop me?!?!?
Loop=7
Loop=8
Loop=9
Exiting now.
```

## Recommendations:

- avoid signals as an IPC mechanism if you can
- especially in multi-threaded programs
- use signals only if you 'have to', eg. for background processes
- more info on signals & C++

<https://en.cppreference.com/w/cpp/utility/program/signal>

# swap ( ) in C and C++

```
/* swap in C
 * pointers are ...
 */
```

```
void
swap(int *x, int *y)
{
    int t = *x;
    *x = *y;
    *y = t;
}
```

```
int a, b;
swap( &a, &b);
```

```
// swap in C++
// references are cool
```


```
void
swap(int &x, int &y)
{
    int t = x;
    x = y;
    y = t;
}
```

```
int a, b;
swap( a, b);
```

```
// swap in C++
// templates are cool
```

```
template <class T>
void swap(T &x, T &y)
{
    T t(x);
    x = y;
    y = t;
}
```

```
double a, b;
swap( a, b);
std::vector<int> c, d;
swap( c, d);
```





# Questions?

# Queues

